



**“Deploy**



**Hugging Face Transformers on AWS  
in just a few lines of code”**



*Nicola Procopio*  
30/05/2023



# About me



**Nicola Procopio**

Senior Data Scientist

## Contacts



[nicola.procopio@finconsgroup.com](mailto:nicola.procopio@finconsgroup.com)



<https://it.linkedin.com/in/nicolaprocopio>



<https://github.com/nickprock>



<https://huggingface.co/nickprock>

Actually @



## Background



## Community



# What's Hugging Face?

*“We are on a mission to democratize **good** machine learning, one commit at a time”*

- **Hugging Face** is the collaboration platform for the machine learning community.
- **The Hugging Face Hub** works as a central place where anyone can share, explore, discover, and experiment with open-source ML.
- With the fast-growing community, some of the most used **open-source ML libraries** and tools, and a talented science team exploring the edge of tech, **Hugging Face is at the heart of the AI revolution.**

★ Models: > 210K  
★ Datasets: > 36K

★ Transformers: 102K stars on GitHub  
★ Datasets: > 16.3K  
★ Diffusers: > 14K  
★ ...

# Hub & Libraries

**Hugging Face** Search models, datasets, users...

Models Datasets Spaces Docs Solutions Pricing

Tasks Libraries Datasets Languages Licenses Other

Filter Tasks by name

Multimodal

- Feature Extraction
- Text-to-Image
- Image-to-Text
- Text-to-Video
- Visual Question Answering
- Document Question Answering
- Graph Machine Learning

Computer Vision

- Depth Estimation
- Image Classification
- Object Detection
- Image Segmentation
- Image-to-Image
- Unconditional Image Generation
- Video Classification
- Zero-Shot Image Classification

Natural Language Processing

- Text Classification
- Token Classification
- Table Question Answering
- Question Answering
- Zero-Shot Classification
- Translation
- Summarization
- Conversational
- Text Generation
- Text2Text Generation
- Fill-Mask
- Sentence Similarity

Audio

- Text-to-Speech
- Automatic Speech Recognition

Models 210,421 Filter by name

new Full-text search 11

**facebook/dino-vitb16**  
Updated 3 days ago • 117M • 44

**bert-base-uncased**  
Updated Nov 16, 2022 • 61.6M • 857

**jonatasgrosman/wav2vec2-large-xlsr-53-english**  
Updated Mar 25 • 60.1M • 126

**xlm-roberta-base**  
Updated Apr 7 • 25.6M • 305

**facebook/dino-vits8**  
Updated 3 days ago • 24.3M • 7

**microsoft/resnet-50**  
Updated Mar 10 • 23.4M • 77

**gpt2**  
Updated Dec 16, 2022 • 20.8M • 1,08k

**facebook/convnext-large-224**  
Updated Mar 2, 2022 • 20.8M • 15

**microsoft/resnet-18**  
Updated 17 days ago • 19.1M • 18

**deepset/sentence\_bert**  
Updated May 19, 2021 • 13.5M • 10

**roberta-base**  
Updated Mar 6 • 11M • 166

**xlm-roberta-large**  
Updated Apr 6 • 9.78M • 143

**bert-base-cased**  
Updated Nov 16, 2022 • 8.6M • 109

## How to use

You can use this model directly with a pipeline for masked language modeling:

```
>>> from transformers import pipeline
>>> unmasker = pipeline('fill-mask', model='bert-base-uncased')
>>> unmasker("Hello I'm a [MASK] model.")
```

# Why Hugging Face on AWS?

**Hugging Face** Deep Learning Containers are available on AWS Sagemaker because there are many advantages:

- **One command is all you need**
- **Accelerate machine learning from science to production (AWS Toolkit)**
- **Built-in performance**
- **Costs**

On Sagemaker you can:

- train
- deploy
- accelerate
- ...

your models

```
from sagemaker.huggingface import HuggingFaceModel

# create Hugging Face Model Class and deploy it as SageMaker Endpoint
huggingface_model = HuggingFaceModel(...).deploy()
```

# Focus on Deploy

Deploy an Hugging Face model is easy with Inference Toolkit.  
First step install and setup the environment.

```
pip install sagemaker --upgrade
```

## Sagemaker Environment

```
import sagemaker  
sess = sagemaker.Session()  
role = sagemaker.get_execution_role()
```

## Local Environment

```
import sagemaker  
import boto3  
  
iam_client = boto3.client('iam')  
role = iam_client.get_role(RoleName='role-name-of-your-iam-role-with-right-permissions')['Role']['Arn']  
sess = sagemaker.Session()
```

# Deploy from Model Hub

- Choose your model from Hugging Face Hub and create a model config.
  - {HF\_MODEL\_ID: checkpoint, HF\_TASK: text-classification}

- Create your model class

```
# create Hugging Face Model Class
huggingface_model = HuggingFaceModel(
    env=hub,
    role=role,
    transformers_version="4.26",
    pytorch_version="1.13",
    py_version='py39',
)
```

- Deploy the endpoint

```
# deploy model to SageMaker Inference
predictor = huggingface_model.deploy(
    initial_instance_count=1,
    instance_type="ml.m5.xlarge"
)
```

**REMEMBER!**  
Each HF\_TASK has an inputs  
format!!!

# Deploy from S3

- Create a *.tar.gz* file contains your model
- Create your role and model class

```
import sagemaker
import boto3

try:
    role = sagemaker.get_execution_role()
except ValueError:
    iam = boto3.client('iam')
    role = iam.get_role(RoleName='sagemaker_execution_role')['Role']['Arn']

print(f"sagemaker role arn: {role}")
```

```
from sagemaker.huggingface import HuggingFaceModel

# create Hugging Face Model Class
huggingface_model = HuggingFaceModel(
    model_data="s3://hf-sagemaker-inference/model.tar.gz", # path to your trained sagemaker model
    role=role, # iam role with permissions to create an Endpoint
    transformers_version="4.26", # transformers version used
    pytorch_version="1.13", # pytorch version used
    py_version="py39", # python version of the DLC
)
```



# Serverless Inference

- Deploy and scale ML models in a Serverless fashion
- Serverless endpoints automatically launch compute resources and scale them in and out depending on traffic similar to AWS Lambda.
- Serverless Inference is ideal for workloads which have idle periods between traffic spurts and can tolerate cold starts.
- Pay-per-use model
- Some Limitations:
  - Memory size
  - Concurrent Invocations
  - Cold Start



# Serverless Inference

## How the code changes?

Only import a module,  
the object initialization and  
a few parameters settings

```
from sagemaker.huggingface.model import HuggingFaceModel
from sagemaker.serverless import ServerlessInferenceConfig

# Hub Model configuration. <https://huggingface.co/models>
hub = {
    'HF_MODEL_ID': 'distilbert-base-uncased-finetuned-sst-2-english',
    'HF_TASK': 'text-classification'
}

# create Hugging Face Model Class
huggingface_model = HuggingFaceModel(
    env=hub,                    # configuration for loading model from Hub
    role=role,                  # iam role with permissions to create an Endpoint
    transformers_version="4.26", # transformers version used
    pytorch_version="1.13",      # pytorch version used
    py_version='py39',          # python version used
)

# Specify MemorySizeInMB and MaxConcurrency in the serverless config object
serverless_config = ServerlessInferenceConfig(
    memory_size_in_mb=4096, max_concurrency=10,
)

# deploy the endpoint endpoint
predictor = huggingface_model.deploy(
    serverless_inference_config=serverless_config
)
```

# Deploy Custom Models

- You need to use libraries not directly supported by the toolkit (es. **SetFit**, Sentence Transformers, SpanMarker, ...)
- Your output format isn't the default
- Other reasons why you cannot use the inference endpoint as it is



# Deploy Custom Models

In our model folder we must add:

1. requirements.txt
2. handler.py
3. code/inference.py

Create a deploy.py, the script which is creating your endpoint.

**requirements.txt** is the file that contains the dependencies, in my case it had only one line: ***setfit***.

For S3 endpoint put it into code folder.

# Deploy Custom Models

***handler.py*** defines:

- how the endpoint must **load the model**
- what the **input** looks like
- what should return as **output**

```
from typing import Dict, List, Any
from setfit import SetFitModel

class EndpointHandler:
    def __init__(self, path=""):
        # load model
        self.model = SetFitModel.from_pretrained(path)
        # ag_news id to label mapping
        self.id2label = {0: "World", 1: "Sports", 2: "Business", 3: "Sci/Tech"}

    def __call__(self, data: Dict[str, Any]) -> List[Dict[str, Any]]:
        """
        data args:
            inputs (:obj: `str`)
        Return:
            A :obj:`list` | `dict`: will be serialized and returned
        """
        # get inputs
        inputs = data.pop("inputs", data)
        if isinstance(inputs, str):
            inputs = [inputs]

        # run normal prediction
        scores = self.model.predict_proba(inputs)[0]

        return [{"label": self.id2label[i], "score": score.item()} for i, score in enumerate(scores)]
```

# Deploy Custom Models

The **code folder** contains ***inference.py***:

- model\_fn

```
def model_fn(model_dir):
    model = SetFitModel.from_pretrained(model_dir)
    return model
```
- input\_fn

```
def input_fn(input_data, content_type):
    decoded_input_data = decoder_encoder.decode(input_data, content_type)
    return decoded_input_data
```
- predict\_fn

```
def predict_fn(data, model):
    setFit_model = model

    inputs = data.pop("inputs", data)
    if isinstance(inputs, str):
        inputs = [inputs]

    # run normal prediction
    scores = model.predict_proba(inputs)[0]

    return [{"label": model_id2label[i], "score": round(score.item(),4)} for i, score in enumerate(scores)]
```
- output\_fn → decode the result

# Deploy Custom Models

Create the endpoint, file ***deploy.py***:

```
hf_model = HuggingFaceModel(  
    model_data="s3://call-your-model/model1.tar.gz",  
    role=role,  
    transformers_version="4.6.1",  
    pytorch_version="1.7.1",  
    source_dir="code",  
    py_version="py36",  
    source_dir="model1/code",  
    entry_point="inference.py",  
)  
  
def predict(input):  
    predictor = HuggingFacePredictor(  
        endpoint_name="huggingface-endpoint")  
    resp = predictor.predict({"inputs": input})  
    print(resp)  
    out = [data['label'] for data in resp if data['score'] >= 0.5]  
    return out  
  
predictor = huggingface_model.deploy(  
    serverless_inference_config = ServerlessInferenceConfig(memory_size_in_mb=6144),  
)
```

# Notes

- **Code repository:**
  - [Hugging Face Sagemaker workshop series](#)
  - [Hugging Face notebooks on Sagemaker](#)
- **Follow Philipp Schmid (Hugging Face ML Eng & AWS ML Hero):**
  - [github](#)
  - [website](#)
- **Blog:**
  - [The Partnership: Amazon SageMaker and Hugging Face](#)
  - [Deploy Hugging Face models easily with Amazon SageMaker](#)





# Thanks!

## LET'S CHAT!

Remember to join  **Hugging Face**